

UNIMORE
UNIVERSITÀ DEGLI STUDI DI
MODENA E REGGIO EMILIA

CMix-NN: Mixed Low-Precision CNN Library for Memory-Constrained Edge Devices

Alessandro Capotondi*, Manuele Rusci, Marco Fariselli,
Luca Benini

**alessandro.capotondi@unimore.it*

Hipert Lab

Department of Physics, Mathematics
and
Informatics

Università di Modena e Reggio Emilia

*Energy-Efficient Embedded Systems
Laboratory*


Dipartimento di Ingegneria dell'Energia
Elettrica e dell'Informazione

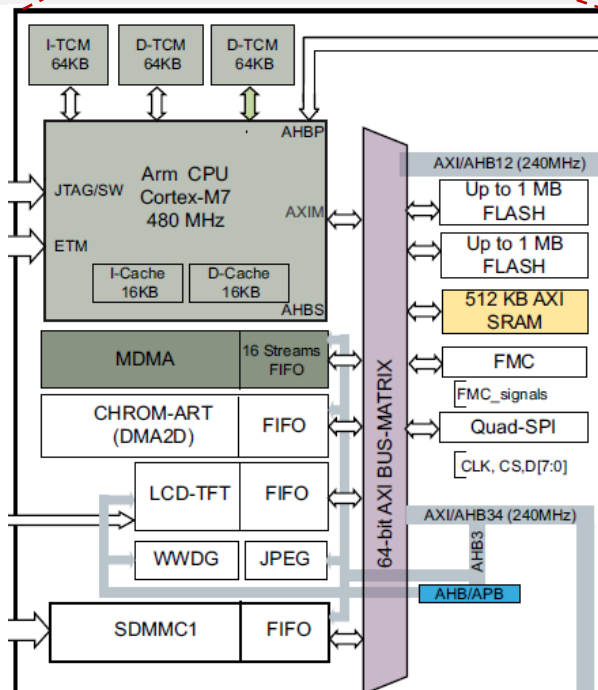
Università di Bologna

ALMA MATER STUDIORUM - UNIVERSITÀ DI BOLOGNA

IL PRESENTE MATERIALE È RISERVATO AL PERSONALE DELL'UNIVERSITÀ DI BOLOGNA E NON PUÒ ESSERE UTILIZZATO AI TERMINI DI LEGGE DA ALTRE PERSONE O PER FINI NON ISTITUZIONALI

Edge Smart Devices






Source: STM32H7 datasheet

- ❑ Low-power (<10-100mW) & low-cost
 - ❑ Smart device are battery- operated
- ❑ Highly-flexible (SW programmable)
- ❑ But **limited resources(!)**
 - ❑ few MB of memories
 - ❑ single RISC core up to few 100s MHz (STM32H7: 480MHz) with DSP SIMD instructions and optional FPU
- ❑ Currently, only tinyML tasks on MCUs (visual wake words, CIFAR10)

Challenge: Run 'complex' and 'big' (Imagenet-size) DL inference on MCU ?

Deep Learning Model Size



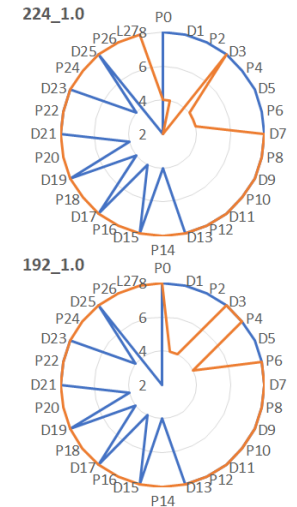
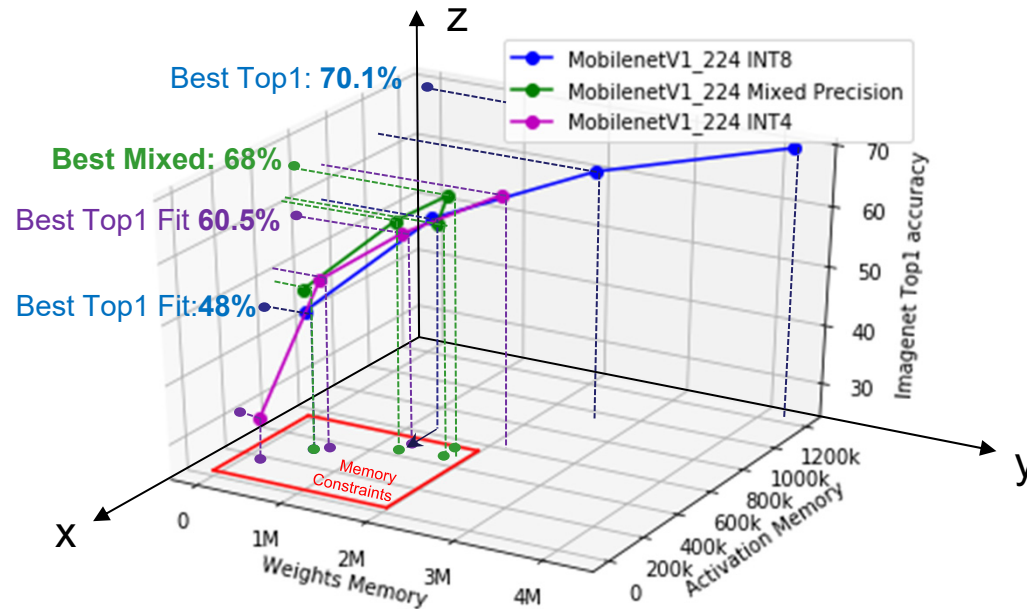
Rusci, Manuele, et al. "Leveraging Automated Mixed-Low-Precision Quantization for tiny edge microcontrollers." arXiv preprint arXiv:2008.05124 (2020).

Fact 1: Integer-only model needed for deployment on low-power MCUs

Fact 2: 8-16 bit are **not** sufficient to bring 'complex' models on MCUs (memory!!)

Mixed-Precision Quantization


Using less than
8 bits...



HAQ (Wang et al 2018)	Accuracy-aware mixed-precision policy search with RL agent. Relies on non-uniform quantization, weight-only.
HAWQ (Dong et al 2019)	Hessian-based metric to compute the quantization policy. Less performant than HAQ.
(Rusci et al 2020)	Mixed Precision Quantization applying minimum tensor-wise quantization ≤ 8 bit to fit the memory constraints


(Wang et al 2018) Wang, Kuan, et al. "Haq: Hardware-aware automated quantization with mixed precision." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2019.
(Dong et al 2019) Dong, Zhen, et al. "Hawq: Hessian aware quantization of neural networks with mixed-precision." *Proceedings of the IEEE International Conference on Computer Vision*. 2019.
(Rusci et al. 2020) "Memory-driven mixed low precision quantization for enabling deep network inference on microcontrollers." *arXiv preprint arXiv:1905.13082* (2019).

Fact 3: Mixed Low-Precision Improves Accuracy on memory constrained device



CMix-NN: Mixed Low-Precision CNN Library

- ✓ **Mixed-precision sub-byte** (2, 4, 8 bit) quantization for weights and activations
- ✓ **Per-Layer** (*PL*), and **Per-Channel** (*PC*) Quantization
- ✓ **Optimized** for ARM ARMv7-M ISA
- ✓ **Open-Source:** <https://github.com/EEESlab/CMix-NN>



CMix-NN Source Code Structure

- Based on CMSIS-NN
 - Same computational model (im2col + GEMM)
 - Same HWC tensor layout
- C Library (not C++)
 - NN Graph is compiled! Good for Latency!
- Convolutional Layer Supported
 - 2D Conv, Depthwise
 - FC

**More than 81 different
convolutional kernel
configurations!**




Cmix-NN final source code is
generated using templates
(mako, python)

Uniform Integer Quantization Flavors

real value
tensor or sub-
tensor $\rightarrow t = S_t \times (T_q - Z_t)$

quantized tensor (INT-Q)

$$S_t = \frac{b - a}{2^Q - 1}$$



$$T_q = Z_t + \text{round} \left(\frac{\text{clamp}(t, a, b)}{S_t} \right)$$

a, b are learned or deterministic params
In general $a \neq b$, **asymmetric quantization**


activations weights

Per-Layer (PL) : asym $[a, b]$ of the whole weight tensor w
Per-Channel (PC) : asym $[a, b]$ of the weight sub-tensor along the outer dim $w[i]$

Typically asymmetric by definition
(e.g. ReLU), b is learned via PACT

$$Y_q = \text{quant}_{act}(x) = Z_y + \text{floor} \left(\frac{\text{clamp}(x, 0, b)}{S_x} \right) \cdot S_x$$

Fake-Quantized Sub-Graph



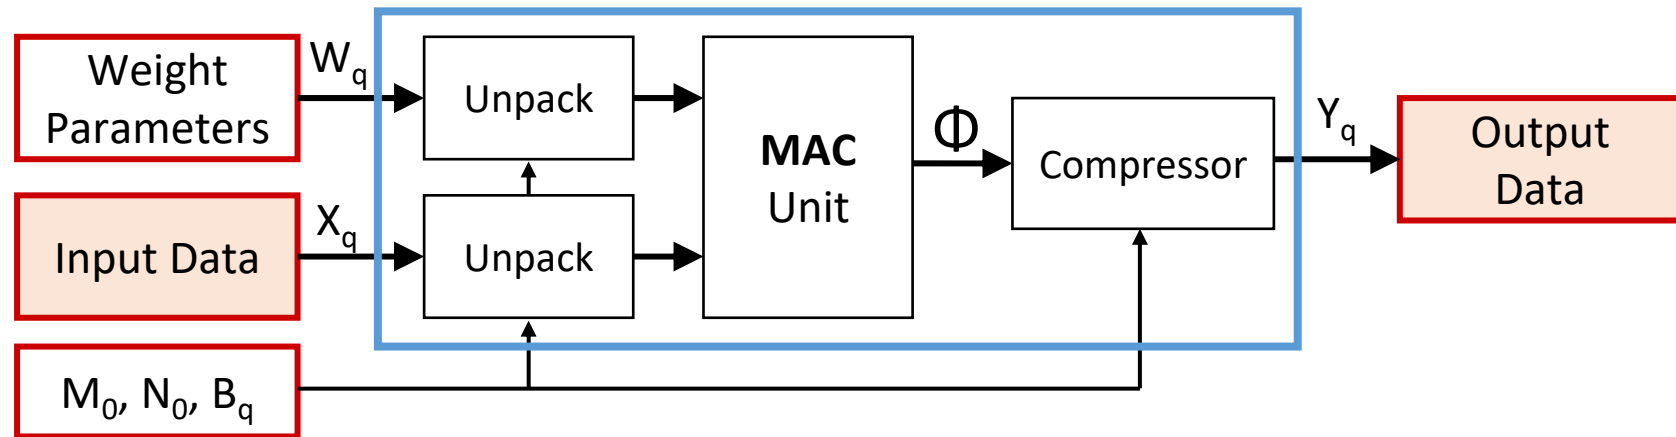
To turn a fake-quantized (sub-byte) sub-graph into an integer-only Representation we use **Integer Channel-Normalization** (Rusci et al. 2019) :

$$Y_q = Z_y + \text{clamp}(\text{floor}(M_0 2^{N_0} (\Phi + B_q)), 0, 2^{Q-1})$$

M_0, N_0, B_q are channel-wise integer params

(Rusci et al. 2019) "Memory-driven mixed low precision quantization for enabling deep network inference on microcontrollers." arXiv preprint arXiv:1905.13082
(Jacob, 2018) Jacob, Benoit, et al. "Quantization and training of neural networks for efficient integer-arithmetic-only inference." CVPR 2018

Layer-by-layer execution



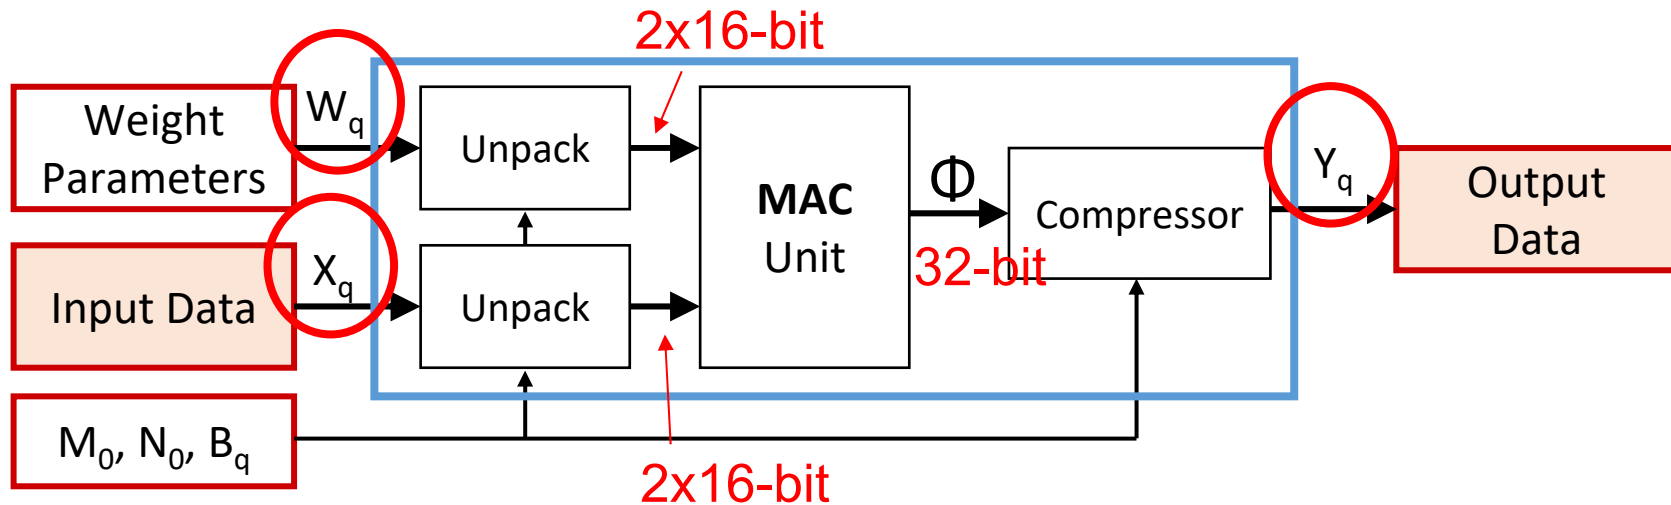
Layer
transfer
function

$$Y_q = Z_y + \text{clamp}(\text{floor}(M_0 2^{N_0} (\Phi + B_q)), 0, 2^{Q-1})$$

$$\Phi = \sum (W_q - Z_w) \cdot (X_q - Z_x)$$

NB. This formulation holds either for PL or PC quantization of weights, only Z_w differs from scalar to array

CMix-NN Quantized Convolutional Layer (2)



- W_q , X_q and Y_q has *arbitrary* precision (2-, 4-, 8-bit)
- MAC Unit exploits *2x16-bit* SIMD MAC (ARMv7-M ISA)
- Compression (or *Activation Function*) brings 32-bit accumulator to *arbitrary* precision (2-, 4-, 8-bit)



CMix-NN Computational Scheme

for $x,y=0; x,y<dimOut; x,y++$ **do**

end for

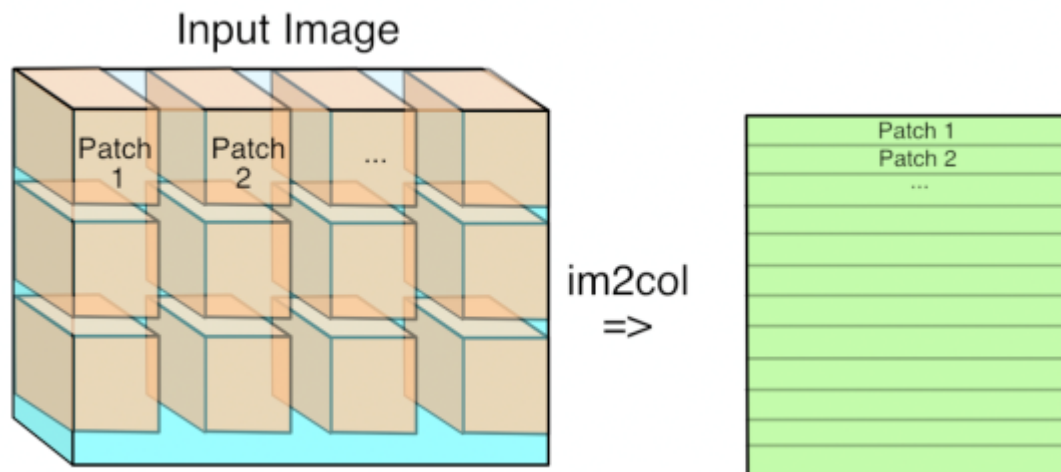
CMix-NN Computational Scheme

im2col

for $x,y=0; x,y<dimOut; x,y++$ do

u32 ptrBuff0 \leftarrow im2colQ(tensorIn[x,y], Zin, kw, kh, InFeat, nCol)

u32 ptrBuff1 \leftarrow im2colQ(tensorIn[x+1,y], Zin, kw, kh, InFeat, nCol)



NOTE! Patches are also packed from original bitwidth to 2x16-bit Vectors

end for

Figure: <https://petewarden.com/2015/04/20/why-gemm-is-at-the-heart-of-deep-learning/>

CMix-NN Computational Scheme

```
for x,y=0; x,y<dimOut; x,y++ do
```

```
  u32 ptrBuff0 ← im2colQ(tensorIn[x,y], Zin, kw, kh, InFeat, nCol)
```

```
  u32 ptrBuff1 ← im2colQ(tensorIn[x+1,y], Zin, kw, kh, InFeat, nCol)
```


```
  for each co in OutFeat; co+=2 do
```

```
    end for
  end for
```


im2col

MatMul

Input Matrix

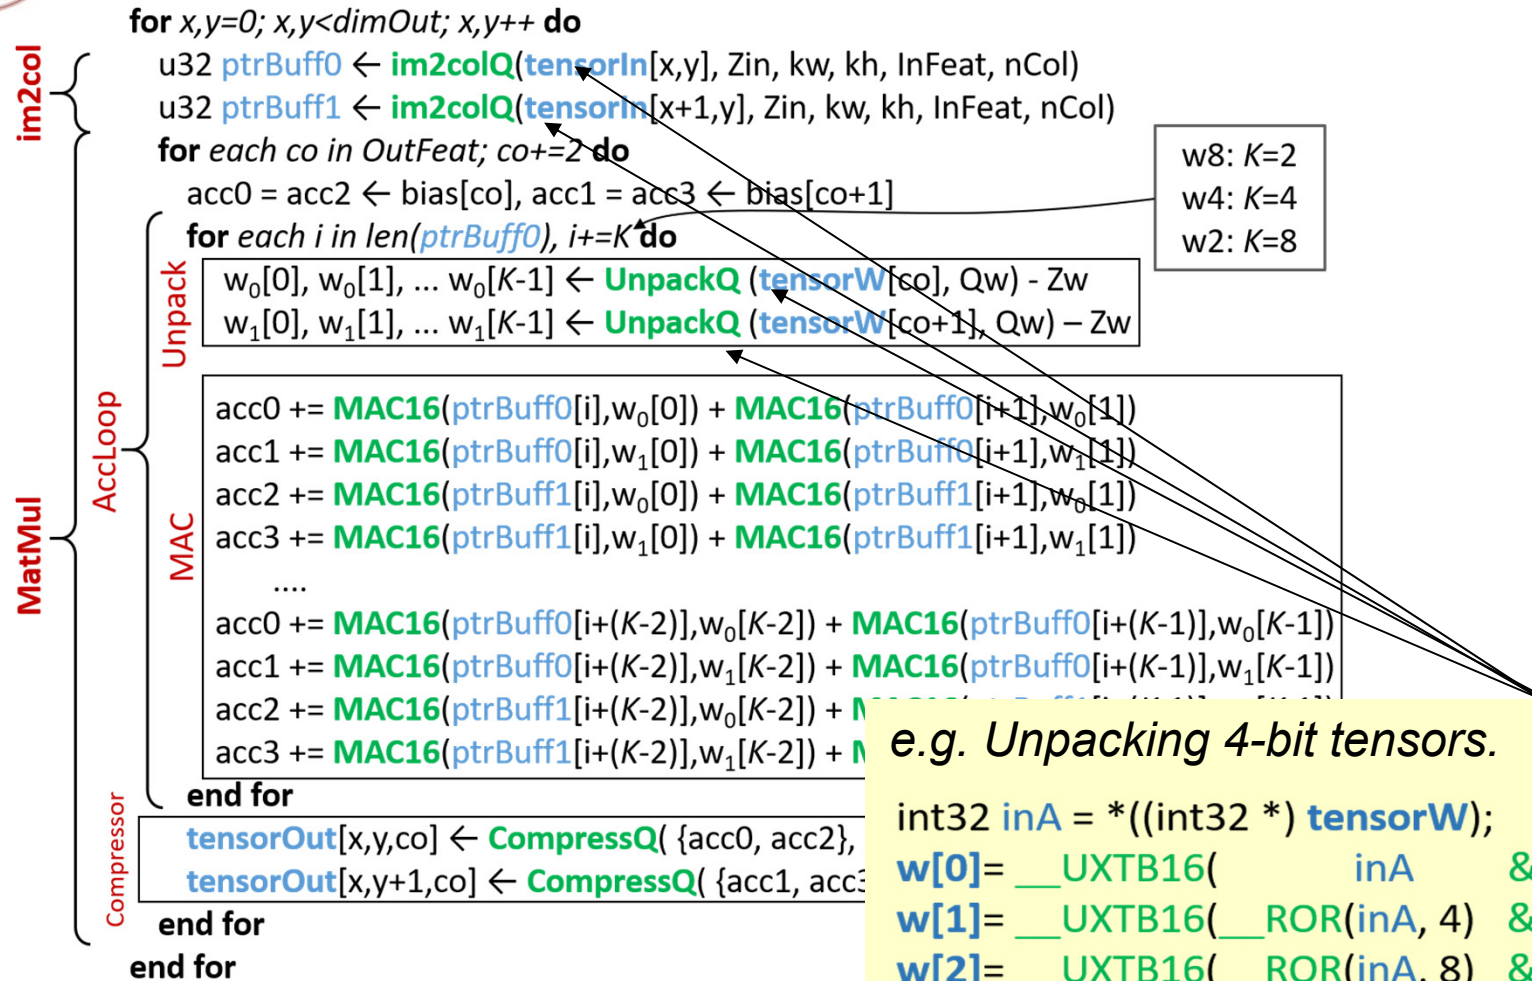


Kernel Matrix




X
K

CMix-NN Computational Scheme



Low-bitdwith Extraction is not a free meal...



Experimental Setup

- **Low-bitwidth Performance Characterization**
- **End2End Network Inference SoA Comparison**

Target Board

- STM32H7
 - ARM Cortex-M7 480 MHz, 2MB FLASH, 512kB SRAM

CNN Network Configuration


- Mobilenet v1 and Mobilenet v2 Family
- Imagenet (1000-class) Classification Task

Inference Libraries SOA

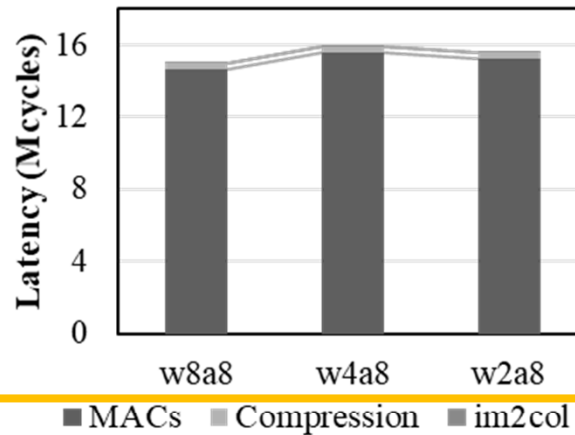
- CMSIS-NN
- STM CubeMX.AI

Low-bitwidth Performance Characterization

Layer 3 – Mobilenet 192_0.5




Layer 27 – Mobilenet 192_0.5



- 👍 Compression and im2col are insensitive
- 👎 MAC loop up to 14% slower (due to unpack)

Varying W bitwidth



- 👍 Compression and MAC are insensitive
 - 👎 im2col loop up to 10% slower
- Varying Act bitwidth

↗ Input dimension = ↘ load/MAC (for this benchmark)
 ↘ W bits = ↗ unpacking instructions

The spatial stationary computation model makes convolution efficiency independent by activation bit

End2End SoA Comparison

	X-CUBE-AI		CMSIS-NN	PC+ICN		CMix-NN (our)				
Network	Mob-V1	Mob-V1	Mob-V1	Mob-V1	Mob-V1	Mob-V1	Mob-V1	Mob-V1	Mob-V1	Mob-V2
Quantization	128_0.25	192_0.5	192_0.5	224_0.75	224_0.75	224_0.75	192_0.5	160_0.50	128_0.25	224_1.0
Tensor Type	—	PL	PL	PC+ICN	PL+ICN	PL+ICN	PC+ICN	PC+ICN	PC+ICN	PC+ICN
	FP32	INT8 ¹	INT8 ²	w4a4	mixed		w8a8	w8a8	w8a8	mixed
Avg Act. Bits	32	8	8	4	6.72	6.72	8	8	8	5.35
Avg Wgt. Bits	32	8	8	4	5.99	5.99	8	8	8	4.61
Model Size (MB)	1.88	1.37	1.37	0.98	1.97	1.97	1.37	1.37	0.49	1.95
Latency (s)	0.206	0.437	0.510	2.290	1.86	1.419	0.677	0.460	0.110	2.013
MAC/Cycle	0.14	0.52	0.45	0.30	0.36	0.48	0.33	0.35	0.26	0.30
Accuracy Top-1	45.0%	59.5%	59.5% ³	60.5%	68.2%	67.0%	62.9%	61.25%	44.6%	54.0%
Energy (μJ)	54.40	115.40	134.64	604.69	491.15	374.70	178.77	121.46	29.05	531.55

Best Cmix-NN solution (MobV1-224-0.75, PC+ICN, Mixed, Top1% 68.2%) up to

- **+23% Accuracy** compared to the best FP32 solution (CubeMX.AI)
- **+9% Accuracy** compared to the best INT8 solution (CubeMX.AI, CMISIS-NN)

Per-Channel quantization is always more accurate compare to Per-Layer (up to **+1.75%**)

Per-Channel quatization is approx. 24% less efficient in MAC/s (CubeMX.AI 0.52 MAC/s, Cmix-NN PL+ICN 0.48, Cmix-NN PC+ICN 0.36)

But, **MobV1-160-0.50**, PC+ICN, has the same or lower latency of any other competitors and provide a **more accurate solution!**


Conclusions

- CMix-NN, flexible inference library for MCU-based edge devices
 - optimized backend to deploy state-of-the-art mixed low-precision deep neural networks for ARM Cortex-M MCUs
 - Support different flavor of quantization
 - Different Datatypes, for weights and activations
 - Per-Channel and Per-Layer quantization
- Open-Source <https://github.com/EEESlab/CMix-NN>
- CMix-NN enables the deployment of a 68% Imagenet MobilenetV1 into a MCU with 2MB FLASH and 512 kB RAM.

What's next

Supporting Complex ML problems on MCU-class require

- Optimized byte and sub-byte operations on ISA!
 - A. Garofalo “XpulpNN: Accelerating Quantized Neural Networks On RISC-V Processors Through ISA Extensions”
 - ARMv8-M (MAC 8/16-bit SIMD arithmetic)
- Novel Quantization Mechanisms to improve accuracy degradation on low-precision (sub-byte) models
- Quantization is only one of the levels for DL Compression
 - Pruning
 - NAS and AutoML
 - How integrate them properly!



Check it out and Contribute to **Cmix-NN!**

<https://github.com/EEESlab/CMix-NN>

Alessandro Capotondi
Università di Modena e Reggio Emilia
alessandro.capotondi@unimore.it

